

Thread Pool Basic Implementation Solutions

Overview

- In the previous lecture, we said that a thread pool needs:
 - A fixed-size container of C++ thread objects
 - A queue of tasks, which are represented as callable objects
- How does our thread pool class implement these requirements?
 - It has `std::vector<std::thread>` as a data member, for the threads
 - It uses the concurrent queue from an earlier lecture
 - The tasks are represented as `std::function` objects
 - The concurrent queue stores these `std::function` objects

Interface

- Describe the interface of the thread pool class
 - The thread pool class has a public submit() member function
 - Users call this to send a task to the thread pool
 - The thread pool class also has a private worker() member function
 - This is the entry point function for the threads
 - The class also has a constructor and a destructor

Member functions

- Describe the member functions of the thread pool class
 - `submit()` takes an `std::function` object as argument (a task)
 - It pushes the `std::function` object onto the concurrent queue
 - The next thread which becomes available will invoke the task
 - `worker()` is an infinite loop
 - It pops a task off the queue, then invokes it
 - When the task is complete, it pops another task off the queue, and so on
 - The constructor populates the `std::vector` of `std::thread` objects
 - These `std::thread` objects have `worker()` as their entry point function
 - The destructor calls `join()` on the `std::thread` objects

thread_pool Usage

- How does other code use the thread pool class?
 - The code creates an object of the thread pool class
 - It then sends tasks to the thread pool
 - To do this, it calls the submit() member function
 - It passes the task as the argument to submit()
 - The task function must have the same signature as the elements of the concurrent queue